

# Annex

## Proof, not paperwork.

Annex reads your codebase, decides where it falls under the EU AI Act, and proves every obligation with a file, a line and a hash.

---

Shivam Gupta · [github.com/shi1720/LexHack](https://github.com/shi1720/LexHack)

#### THE PROBLEM

## Every AI Act dossier is a document a company wrote about itself.

Nobody has ever checked one against the system it describes. That is not a snipe at compliance teams — it is structural. The Act asks whether your system logs inference, whether a human can override it, whether you examined your data for bias. **Those answers are in the repository.** The people who write the dossier cannot read it, and the people who can were never asked.

~90

#### AI GOVERNANCE VENDORS

In the IAPP's 2026 category census.

2

#### MENTION SOURCE CODE

One is a GDPR privacy scanner. The other is the phrase "a U.S. codebase".

40%

#### CAN'T BE CLASSIFIED

Of 106 enterprise AI systems studied by appliedAI, the risk tier was unclear.

# The deadline moved. The obligations did not.

On 27 July 2026 the Digital Omnibus pushed the Annex III high-risk deadline from August 2026 out to **2 December 2027**, and the compliance industry exhaled. **It left Article 50 exactly where it was.**

## In force today

|                   |                                            |
|-------------------|--------------------------------------------|
| <b>2025-02-02</b> | Article 5 prohibited practices · €35M / 7% |
| <b>2025-08-02</b> | GPAI model obligations                     |
| <b>2026-08-02</b> | <b>Article 50 transparency · €15M / 3%</b> |
| <b>2023-07-05</b> | NYC Local Law 144 bias audits              |

## What everyone is waiting for

|                   |                                           |
|-------------------|-------------------------------------------|
| <b>2026-12-02</b> | Art. 50(2) marking, grandfathered systems |
| <b>2026-12-02</b> | New Art. 5 prohibitions (NCII, CSAM)      |
| <b>2027-12-02</b> | Annex III high-risk — moved from Aug 2026 |
| <b>2028-08-02</b> | Annex I embedded high-risk                |

If your product talks to a person or generates content, **you have been in scope since August**. The deadline nobody is talking about is the one that already passed.

## WHAT IT DOES

# One scan. Four artefacts. 220 milliseconds.

### PROHIBITED

Article 5(1)(f) – emotion inference in the workplace

classification Annex III, point 4(a) · recruitment and candidate selection 97% confidence  
src/screening/rank.ts:28

```
const decision = candidateScore >= ADVANCE_THRESHOLD ? 'advance' : 'reject';
```

- \* missing Art. 5(1)(f) **IN FORCE** emotion inference from video interviews
- \* missing Art. 50(1) **IN FORCE** the chat UI never says it is an AI
- \* missing Art. 14 human oversight – no review step, no override, no stop control
- \* missing LL 144 §5-301 no independent bias audit in the last twelve months

conformity  4/100 in force today 4/100  
exposure €294,000 statutory ceiling ledger **53CC-FAD4-DDAF-DD63** 14 files · 47 ms

### Evidence explorer

Statute on the left, the lines that answer it on the right.

### Annex IV dossier

Headings in EN · DE · FR, every unanswerable point left open.

### A pull request

Additive files that close what code can close.

### Standards exports

SARIF, CycloneDX attestation, ML-BOM.

## Not a paperwork problem. A verification problem.

A frontier model can draft a plausible Annex IV document this afternoon. What it cannot do is tell you whether the document is *true*. So we built for verification, and two things follow from that.

### The law gets a test suite

A control is a function, so it can be tested. 49 golden fixtures across all five packs — each a minimal repository and the status the control must return — plus a suite that checks every control against a written-down Article 113 table. If a regex gets greedier, the obligation fails the build rather than silently mis-reporting somebody's conformity.

### The evidence ledger

Every result is hashed into a chain: the control, its status, the rule-pack version, the SHA-256 of every file it cites. The root goes on the dossier's front page. Re-run on the same commit and you get the same root. `annex verify` re-derives every entry from the results the report describes, so an edited status stops matching, and `--against` re-hashes each cited file off disk. A checksum chain, not a signature — it makes a silent edit detectable, not impossible.

That is the whole difference between a document and a proof.

## Substantial modification, caught in the pull request.

Article 3(49): a change that affects compliance with Chapter III is a **substantial modification**. Article 43(4): that re-opens the entire conformity assessment. It is triggered by a code change — so only the tool that watches code changes can see it.

```
$ annex diff --base origin/main --head HEAD
```

**SUBSTANTIAL MODIFICATION** origin/main → HEAD

The risk classification moved from high to prohibited. Under Article 25(1)(c) and Article 43(4) this is a substantial modification: the conformity assessment has to be re-opened and the technical documentation updated.


```
conformity 81 → 2 (-79)
```

```
✖ Effective human oversight while the system is in use satisfied → missing
```

A questionnaire structurally cannot do this. It is also, not coincidentally, the event that makes the revenue recurring.

## HOW IT WORKS

# A regulation is a function over a repository.



### THE TWO THINGS THAT MAKE IT EVIDENCE

- **The law gets a test suite.** 49 golden fixtures across all five packs, so a detector that drifts fails the build.
- **The ledger is checkable.** annex verify re-derives every entry from the results; --against re-hashes each cited file off disk.

## It is a classifier, so it owes you its error rate.

42 hand-labelled cases. Roughly half exist to catch **false positives**: card-fraud detection (excluded from Annex III 5(b)), one-to-one identity verification (excluded from 1(a)), campaign logistics (excluded by 8(b) — nobody sees the output), invoice OCR (Art. 50(2) does not reach re-expression), a mood-journal app that is 1(c) high-risk and *not* the Art. 5(1)(f) prohibition, and a blog post about prohibited practices.

**100%**

RISK-TIER ACCURACY

**100%**

FINDING RECALL

**100%**

CARVE-OUT PRECISION

### A clean sheet is a claim about the corpus, not the world

These labels were written by the people who wrote the detectors. The report says so, and five cases exist to find the edge: a lending decision expressed only in SQL, and a chatbot whose AI-ness may or may not be “obvious to a reasonably well-informed person”.

### Building it found four real bugs

Domain signals firing on prose; a missing Art. 50(2) rule; broken go .mod parsing; an Art. 50(2) carve-out stated in the caveat and never implemented. Then the self-scan found a GDPR Article 17 gap *in us* — which we fixed.

## Three people, one artefact each.

| Who                | What they feel                                                     | What Annex gives them                            |
|--------------------|--------------------------------------------------------------------|--------------------------------------------------|
| Engineering lead   | "Legal keeps asking me questions about our own code."              | A CI check and a pull request. Zero procurement. |
| Founder / CTO      | "An enterprise deal is blocked on an AI governance questionnaire." | A trust page to send back today.                 |
| Legal / compliance | "I am signing a document I cannot verify."                         | Evidence with a citation and a hash.             |

### Pricing follows the statute's unit of account

**Free** — CLI, GitHub Action, SARIF, classification.

**\$35/contributor/mo** — blocking PR check, remediation, drift.

**\$12–25k per governed AI system/yr** — dossier, attestation, freshness.

**\$75–250k** — enterprise, self-hosted, custom detectors.

### The anchors it prices against

€9,900 boutique readiness sprint · \$15–50k consultancy · €75k+ Big-4 programme · €10–40k third-party conformity assessment per system · €193–330k to stand up a QMS, plus €71k/year (EC impact assessment).

# The mapping is a commodity. Four other things are not.

## Defensible

- **False-positive economics.** Telling a real Art. 14 control from an if approved : branch, across every framework. Endor Labs raised \$93M on exactly this problem in security.
- **Acceptance history.** Which code patterns an assessor actually accepts can only be learned from dossiers accepted and rejected.
- **Workflow position.** A required check on the default branch, plus code Annex wrote that now lives in the customer's repo.
- **Drift detection.** Triggered by a code change; only we see code changes.

## What would make me wrong

- The Commission must publish a **free simplified Annex IV form** for SMEs. That is why we sell verification, not generation.
- Code presence is not conformity, and assessors know it. If they treat it as supporting material, the price ceiling drops.
- The category is **\$492M in 2026** (Gartner). This is a wedge into GRC and AppSec, not a standalone category play.
- Budgets follow deadlines. Expect an air pocket until H2 2027.

BUILT FOR THIS HACKATHON

**45 obligations. 5  
jurisdictions. 0 model calls in  
the core.**

**45**

EXECUTABLE OBLIGATIONS

**75**

SIGNAL DETECTORS

**75**

TESTS, INCL. GOLDEN FIXTURES

**220<sub>ms</sub>**

TO SCAN 168 FILES

**Rule packs, each reconciled against primary sources**

|                          |                                     |
|--------------------------|-------------------------------------|
| <b>EU AI Act</b>         | 24 obligations · post-Omnibus dates |
| <b>GDPR</b>              | 5 · Arts. 9, 13–17, 22, 35          |
| <b>NYC Local Law 144</b> | 5 · the arithmetic is in the rule   |
| <b>Colorado ADMT Act</b> | 5 · SB 26-189, from Jan 2027        |
| <b>NIST AI RMF 1.0</b>   | 6 · version-pinned subcategories    |

**Where the language model is — and is not**

The engine never calls one. Classification and control evaluation are rule-based and offline, because an auditor cannot accept "*the model thought so*" as evidence.

There is exactly one model call, labelled `model-written` in the UI: rewriting an **already settled** finding for an engineer, a founder or an assessor. It receives the decision as fact and cannot change it.

# A GitHub App, and a dossier in front of a notified body.

The CLI and SARIF output already work in CI. The App turns a scan anyone can ignore into a gate nobody can. After that, the question that decides whether this is a product or a demo is whether an assessor — TÜViT, Nemko, Bureau Veritas — will accept code-grounded evidence as a starting artefact. That is a question about institutions, not code, and it is answered by knocking on doors.

---

## Proof, not paperwork.

Shivam Gupta · LexHack 2026 · [github.com/shi1720/LexHack](https://github.com/shi1720/LexHack)