

Technical documentation referred to in Article 11(1)

Regulation (EU) 2024/1689, Annex IV · compiled from source by Annex

System	HireFlow
Provider	HireFlow
Version under assessment	3549765d90a3
Issued	2026-09-15
Classification	High-risk under Annex III — employment: recruitment and candidate selection.
Role under Article 3	provider and deployer
Evidence ledger	050D - 30FF - C32B - EA27
Evidence citations	41
Open items	14

CONTENTS

1. General description of the AI system
2. Detailed description of the elements of the AI system and of the process for its development
3. Monitoring, functioning and control of the AI system
4. Appropriateness of the performance metrics
5. Risk management system in accordance with Article 9
6. Relevant changes made through the lifecycle
7. Harmonised standards applied, or solutions adopted
8. Copy of the EU declaration of conformity (Article 47)
9. Post-market performance evaluation system (Article 72)

§ Statement of method

This document was compiled by Annex directly from the source code of the system it describes, at the commit identified above. Every statement below falls into one of three categories, and they are never mixed:

1. **Evidenced.** The statement is supported by one or more citations to a file and a line in the repository. Each citation carries the SHA-256 digest of the file it came from, recorded in Appendix B.
2. **No evidence found.** The engine searched the whole snapshot and found nothing supporting the statement. The negative finding is recorded rather than omitted.

3. **Open.** The determination requires a judgement that cannot be derived from source code — a residual-risk acceptance, a declared accuracy level, an accountable person. These are listed per section and summarised on the front page. They are left blank deliberately: a generated document that invents them would be a false statement to a competent authority, which Article 99(5) prices at EUR 7 500 000 or 1 % of worldwide annual turnover.

No language model participated in any determination in this document. Classification and control evaluation are rule-based, deterministic and offline; the same commit always produces the same evidence ledger root. Where a language model contributed narrative prose, the section says so.

This document is a technical artefact, not legal advice, and it is not a conformity assessment. It is the input a competent person needs in order to carry one out, and the means by which a third party can check that its claims are true.

1 General description of the AI system

(a) Intended purpose, provider and version. Ranks and prioritises job applicants against a posting so recruiters can prioritise review. Produces a recommendation; a person always makes the decision.

Provider: HireFlow. Version under assessment: 3549765d90a3 .

(b) Interaction with other systems. The system calls externally provided models: OpenAI model provider. Under Article 25(1)(c), directing a general-purpose AI system at an Annex III use case makes this repository's owner the provider of a high-risk AI system.

(c) Software versions. 7 production dependencies were resolved from the manifests in this repository. The full inventory is emitted as a CycloneDX ML-BOM alongside this document.

(d) Forms in which the system is placed on the market. Delivered as a hosted software service with a user-facing interface.

(e) Hardware. Not determinable from the codebase. Describe the runtime environment and any hardware requirement.

(f) Product photographs and layout. Not applicable: the system is software, not a component of a physical product.

(g)-(h) Deployer interface and instructions for use. Deployer-facing instructions including a limitations section were found.

EVIDENCE

docs/ai-act/role-determination.md:1

Role under the EU AI Act

README.md:1

HireFlow

docs/ai-act/instructions-for-use.md:1

Instructions for use

docs/ai-act/instructions-for-use.md:11

(i) Intended purpose

docs/ai-act/risk-management.md:36

3. **Inform.** Instructions for use published to deployers; candidate notice

docs/ai-act/instructions-for-use.md:70

Limitations

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Hardware and runtime environment (point 1(e))
- ☐ Forms of placement on the market (point 1(d))

2 Detailed description of the elements of the AI system and of the process for its development

- (a) Development methods and recourse to third-party systems.** The system is built on pre-trained models supplied by third parties (OpenAI model provider). Article 25(2) obliges the initial provider to make available the technical documentation needed to assess Article 16 compliance — unless it has specified that its system is not to be changed into a high-risk AI system, which switches that duty off. Confirm which applies before relying on it.
- (b) Design specifications.** The general logic is expressed in system prompts and orchestration code located in the evidence below. What the system optimises for, and the trade-offs accepted to meet Chapter III Section 2, are design decisions that must be recorded here.
- (c) System architecture and computational resources.** 34 source files were analysed. Not determinable from the codebase: describe how components feed into each other and the resources used to develop, train, test and validate the system.
- (d) Data requirements and datasheets.** Dataset documentation was found.
- (e) Human oversight measures (Article 14).** All three oversight affordances were found: a human review step, an override path, a stop control.
- (f) Pre-determined changes.** Not determinable from the codebase. Record any changes to the system and its performance fixed at the time of the initial conformity assessment.
- (g) Validation and testing procedures.** An evaluation suite and declared accuracy figures were both found. Annex IV point 2(g) also requires metrics for potentially discriminatory impacts, and dated, signed test reports.
- (h) Cybersecurity measures.** Input defences and adversarial testing were both found.

EVIDENCE

evals/bias-audit.ts:4

* NYC Local Law 144, 6 RCNY § 5-301(b): calculate the selection rate and the

evals/bias-audit.ts:23

export function selectionRate(rows: AuditRow[]): number {

src/app/careers/notice.tsx:14

export const BIAS_AUDIT_SUMMARY_URL = '/careers/bias-audit-2026';

src/app/careers/notice.tsx:9

* § 5-303: the bias audit summary must be published on the employment section

docs/ai-act/risk-management.md:42

on 2026-09-02, on the basis of the Q2 monitoring data and the current bias audit.

docs/ai-act/data-governance.md:1

Data and data governance

docs/ai-act/data-governance.md:12

Training data source: 38,000 historical applications from consenting customer

README.md:35

| `data-governance.md` | Article 10 |

src/screening/rank.ts:4

import { gate } from '../lib/ai-act/human-oversight';

src/screening/rank.ts:14

decision: 'advance' | 'pending_review';

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ General logic, key design choices and optimisation target (point 2(b))
- ☐ System architecture and computational resources (point 2(c))
- ☐ Pre-determined changes (point 2(f))
- ☐ Dated and signed test reports (point 2(g))

3 Monitoring, functioning and control of the AI system

Capabilities and limitations in performance. Not determinable from the codebase from source alone. Annex IV point 3 asks specifically for the degrees of accuracy **for specific persons or groups** on whom the system is used, not only the overall figure.

Foreseeable unintended outcomes and sources of risk. Classification identified: Employment: recruitment and candidate selection; Direct interaction with natural persons; Generation of synthetic text. Each is a source of risk to fundamental rights that Article 9 must address.

Human oversight measures. All three oversight affordances were found: a human review step, an override path, a stop control.

Input data specifications. Not determinable from the codebase. Specify the expected input format, range and quality.

EVIDENCE

```
lib/ai-act/audit-log.ts:38
export function recordInference(input: Omit<InferenceRecord, 'decisionId' |
'occurredAt' | 'inputDigest' | 'outputDigest'> & {

lib/ai-act/audit-log.ts:59
console.info(JSON.stringify({ event: 'ai.inference_decision', ...record }));

src/app/api/candidate-data/route.ts:24
// Right to erasure: personal data is deleted, the decision log is pseudonymised.

src/screening/prompts.ts:1
export const PROMPT_VERSION = 'screen-2026-09-02';

lib/ai-act/audit-log.ts:14
modelId: string;

src/screening/rank.ts:4
import { gate } from '../lib/ai-act/human-oversight';
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Accuracy for specific persons or groups
- ☐ Input data specifications

4 Appropriateness of the performance metrics

An evaluation harness was found in the repository (see evidence). Why those metrics are the appropriate ones for this intended purpose is an argument that must be written here — the presence of a metric is not the same as its appropriateness.

EVIDENCE

`evals/accuracy.ts:9`

```
export interface EvalCase { resumeText: string; postingId: string; humanLabel:  
'advance' | 'reject' }
```

`evals/accuracy.ts:6`

```
export const GROUND_TRUTH_SET = 'evals/fixtures/screening-golden-2026Q2.jsonl';
```

`docs/ai-act/instructions-for-use.md:20`

```
| Accuracy | Agreement with human recruiter label | **0.86** | 420-case frozen set,  
'evals/accuracy.ts' |
```

`docs/ai-act/instructions-for-use.md:24`

Accuracy degrades on: resumes under 200 words, non-English resumes, and career

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Justification that the chosen metrics suit the intended purpose

5 Risk management system in accordance with Article 9

Evidenced — Risk management system across the lifecycle (EU AI Act Art. 9). A risk register with an explicit residual-risk judgement was found.

Article 9(5) requires residual risk to be judged acceptable for each individual hazard and overall. That judgement, and the person who made it, cannot be derived from code.

EVIDENCE

docs/ai-act/incident-reporting.md:28

5. Investigate, including a risk assessment and corrective action.

docs/ai-act/post-market-monitoring.md:26

Reviewed quarterly. Feeds section 3 of the risk management system.

docs/ai-act/risk-management.md:1

Risk management system

docs/ai-act/risk-management.md:39

5. Residual risk

docs/ai-act/risk-management.md:41

Per-hazard and overall residual risk judged **acceptable** by the VP Engineering

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Residual risk acceptance and the accountable person (Article 9(5))

6 Relevant changes made through the lifecycle

This document describes the system at commit `3549765d90a3` . Annex maintains change history by comparing consecutive scans: where a change affects compliance with Chapter III Section 2 or modifies the intended purpose, it is a substantial modification within Article 3(49) and reopens the conformity assessment under Article 43(4).

Not determinable from the codebase: no change control artefacts were found. Article 17(1)(a) expects management of modifications to be part of the quality management system.

EVIDENCE

```
tests/rank.test.ts:10
describe('rankApplicants', () => {
```

```
tests/rank.test.ts:11
it('sorts advanced applicants by score', async () => {
```

```
.github/workflows/ai-act-conformity.yml:8
runs-on: ubuntu-latest
```

```
.github/workflows/test.yml:5
runs-on: ubuntu-latest
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Log of relevant changes since the initial conformity assessment

7 Harmonised standards applied, or solutions adopted

No harmonised standards for the AI Act have yet been cited in the Official Journal of the European Union. CEN-CENELEC JTC 21 has prEN 18286 (quality management system for AI Act regulatory purposes) at formal vote. Until a reference is published, Annex IV point 7 takes its second branch: a detailed description of the solutions adopted to meet the Chapter III Section 2 requirements, and a list of other relevant standards applied.

No other standards or technical specifications were referenced in this repository.

The solutions adopted are the controls assessed in this document; each is listed with its evidence in the appendix.

EVIDENCE

```
tests/rank.test.ts:10
describe('rankApplicants', () => {

tests/rank.test.ts:11
it('sorts advanced applicants by score', async () => {

.github/workflows/ai-act-conformity.yml:8
runs-on: ubuntu-latest
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ List of other standards and technical specifications applied

8 Copy of the EU declaration of conformity (Article 47)

Article 47 requires a written, machine-readable, physically or electronically signed EU declaration of conformity for each high-risk AI system, kept available to national competent authorities for ten years.

Not determinable from the codebase: no EU declaration of conformity was found. It cannot be drawn up until the open items in this document are closed — by drawing it up, the provider assumes responsibility for compliance.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ EU declaration of conformity (Article 47, Annex V)

9 Post-market performance evaluation system (Article 72)

Evidenced — Post-market monitoring plan (EU AI Act Art. 72(1)). A post-market monitoring plan was found.

Article 72(3) makes the post-market monitoring plan part of this technical documentation. The Commission template is not due until 2 September 2027, so the plan is free-form until then.

EVIDENCE

```
docs/ai-act/post-market-monitoring.md:1
# Post-market monitoring plan

docs/ai-act/risk-management.md:22
## 3. Risks from post-market monitoring

docs/ai-literacy.md:19
goes into the post-market monitoring review.
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Post-market monitoring plan (Article 72(3))

A Appendix A — Control register

Control	Citation	Status	Finding
Measures to support AI literacy	EU AI Act Art. 4	SATISFIED	Documentation referring to AI training or usage policy was found.
Tell people they are talking to an AI	EU AI Act Art. 50(1)	SATISFIED	An AI disclosure string was found in user-facing interface code.
Mark synthetic output in a machine-readable format	EU AI Act Art. 50(2)	MISSING	This system generates synthetic content and no machine-readable marking was found.
Determine and record whether you are the provider	EU AI Act Art. 3(3)	SATISFIED	The repository records its role under the AI Act in writing.
Risk management system across the lifecycle	EU AI Act Art. 9	SATISFIED	A risk register with an explicit residual-risk judgement was found.
Examine training and evaluation data for bias	EU AI Act Art. 10(2)(f)	SATISFIED	Bias or fairness testing was found in the codebase.

Control	Citation	Status	Finding
Document data provenance and preparation	EU AI Act Art. 10(2)	SATISFIED	Dataset documentation was found.
Annex IV technical documentation exists and is current	EU AI Act Art. 11(1)	MISSING	No Annex IV technical documentation was found.
Automatic logging of events over the system lifetime	EU AI Act Art. 12(1)	SATISFIED	Inference logging was found, carrying both a model version and a traceable identifier.
Retain automatically generated logs for at least six months	EU AI Act Art. 19(1)	SATISFIED	A log retention period meeting the six-month floor was declared.
Instructions for use for deployers	EU AI Act Art. 13(3)	SATISFIED	Deployer-facing instructions including a limitations section were found.
Effective human oversight while the system is in use	EU AI Act Art. 14(1)	SATISFIED	All three oversight affordances were found: a human review step, an override path, a stop control.
Declared accuracy metrics measured against a fixed set	EU AI Act Art. 15(1)	SATISFIED	An evaluation suite and declared accuracy figures were both found.
AI-specific cybersecurity measures	EU AI Act Art. 15(5)	SATISFIED	Input defences and adversarial testing were both found.
Quality management system	EU AI Act Art. 17(1)	PARTIAL	Found a CI pipeline and an automated test suite, which is a start on Article 17.
Post-market monitoring plan	EU AI Act Art. 72(1)	SATISFIED	A post-market monitoring plan was found.
Serious incident reporting procedure	EU AI Act Art. 73(2)	SATISFIED	An incident procedure naming the statutory reporting deadlines was found.
Registration in the EU database	EU AI Act Art. 49(1)	MISSING	No reference to EU database registration was found.
Data protection impact assessment	GDPR Art. 35(1)	MISSING	No data protection impact assessment was found.
Erasure and rectification of personal data	GDPR Art. 16	SATISFIED	Data subject rights endpoints were found.

Control	Citation	Status	Finding
Lawful basis for special-category data	GDPR Art. 9(1)	PARTIAL	Special-category attributes appear in the data model, and a consent or lawful-basis mechanism exists.
Annual independent bias audit	NYC Local Law 144 § 5-301(a)	SATISFIED	Bias audit documentation was found.
Selection rate and impact ratio by category	NYC Local Law 144 § 5-301(b)	SATISFIED	Impact-ratio arithmetic including intersectional categories was found.
Publish the audit summary before use	NYC Local Law 144 § 5-303	SATISFIED	A published bias audit summary was found.
Ten business days notice to candidates	NYC Local Law 144 § 5-304	SATISFIED	Candidate notice including accommodation instructions was found.
Publish the data retention policy and data sources	NYC Local Law 144 § 20-871(b)(3)	SATISFIED	A retention policy and data source documentation were both found.
GOVERN 1.1 — legal and regulatory requirements are understood and documented	NIST AI RMF GOVERN 1.1	SATISFIED	The repository documents the regulatory regimes that apply to it.
GOVERN 1.6 — an inventory of AI systems is maintained	NIST AI RMF GOVERN 1.6	PARTIAL	Model identifiers appear in the code, but there is no single inventory artefact.
MAP 1.1 — intended purpose and deployment context are documented	NIST AI RMF MAP 1.1	SATISFIED	The intended purpose and deployment context are documented.
MEASURE 2.11 — fairness and bias are evaluated and documented	NIST AI RMF MEASURE 2.11	SATISFIED	Fairness evaluation was found in the codebase.
MEASURE 2.7 — security and resilience are evaluated	NIST AI RMF MEASURE 2.7	SATISFIED	Security controls on the model path were found.
MANAGE 4.1 — post-deployment monitoring plans are implemented	NIST AI RMF MANAGE 4.1	SATISFIED	Post-deployment mechanisms found: monitoring, appeal or override, incident response, change management.

B Appendix B — Evidence ledger

Algorithm `sha256-chain/v1`. Root

`050d30ffc32bea27eea98341865576c6453e9a458e8678e8838b78dea29438b6`.

Each entry hashes the previous entry, the control identifier, its status and score, the rule-pack version, and the digest of every piece of evidence cited. Re-running Annex on the same commit reproduces this root exactly. If a cited file changes by one character, the chain breaks and `annex verify` names the entry that stopped matching.

#	Control	Status	Entry hash
0	eu-ai-act.art10.bias-examination	satisfied	ecec7f1935e778498061
1	eu-ai-act.art10.data-governance	satisfied	52a4b5b389ff31f0af8e
2	eu-ai-act.art11.technical-documentation	missing	627aa0acfb020c2dafb4
3	eu-ai-act.art12.record-keeping	satisfied	ec1ebd406db54d140041
4	eu-ai-act.art13.instructions-for-use	satisfied	18fa15d169dc06a4b579
5	eu-ai-act.art14.human-oversight	satisfied	00e1129a89140e18ec04
6	eu-ai-act.art15.accuracy	satisfied	3e0610fefee404c00e6d
7	eu-ai-act.art15.cybersecurity	satisfied	b406573230e5faa30ed0
8	eu-ai-act.art17.quality-management	partial	053756f1296783ba771c
9	eu-ai-act.art19.log-retention	satisfied	41d93dff2062eb8c3375
10	eu-ai-act.art25.role-determination	satisfied	2d4d3b2d050d4aaa97d8
11	eu-ai-act.art4.ai-literacy	satisfied	c1563a4106f741010b49
12	eu-ai-act.art49.registration	missing	02eb597cfd8810a8ea35
13	eu-ai-act.art5.emotion-workplace	not applicable	a15b85ad33930cf7ce55
14	eu-ai-act.art5.face-scraping	not applicable	c5cadbfd6eb86d6a34b5
15	eu-ai-act.art5.ncii-csam-safeguards	not applicable	29953d8381eac2d71ccc
16	eu-ai-act.art5.social-scoring	not applicable	5776d17439d48ce1dc5d
17	eu-ai-act.art50.1.interaction-disclosure	satisfied	20218a493c2ee6ba07c8
18	eu-ai-act.art50.2.content-marking	missing	b7108a6047bb53d119e1
19	eu-ai-act.art50.3.biometric-notification	not applicable	4c5be3a41c46b11d20d8
20	eu-ai-act.art50.4.deepfake-labelling	not applicable	8c24a0cd2a3de6c70968
21	eu-ai-act.art72.post-market-monitoring	satisfied	e9d42682ebdf2109ecc1
22	eu-ai-act.art73.incident-reporting	satisfied	9c3eaeac3c54b1dd5ade
23	eu-ai-act.art9.risk-management	satisfied	de29845c98857d133583
24	gdpr.art13-15.meaningful-information	not applicable	e2fc8c1c1936961742ec

#	Control	Status	Entry hash
25	gdpr.art17.erasure	satisfied	e7cb79c7ee8a27e1015b
26	gdpr.art22.human-intervention	not applicable	90a151b00d6afd302209
27	gdpr.art35.dpia	missing	c0607c930d0a1cc56fc2
28	gdpr.art9.special-category	partial	478741223ffeb84fd40d
29	nist-ai-rmf.govern-1-1	satisfied	a0f37344d104e7715879
30	nist-ai-rmf.govern-1-6	partial	1ca2fad432f9ea0fe061
31	nist-ai-rmf.manage-4-1	satisfied	efaa8cef9480501b272b
32	nist-ai-rmf.map-1-1	satisfied	662b4b7bd8858937aa27
33	nist-ai-rmf.measure-2-11	satisfied	b26ba8904f8145095bfd
34	nist-ai-rmf.measure-2-7	satisfied	9457b8d94968ff75f81d
35	nyc-ll144.bias-audit	satisfied	b0f7d88e28a5f3ac3f5a
36	nyc-ll144.candidate-notice	satisfied	643880c871f7d5e0a02e
37	nyc-ll144.data-policy	satisfied	e41b8f792ec8f5b41b65
38	nyc-ll144.impact-ratio	satisfied	7ce07548bddca488c78f
39	nyc-ll144.publish-summary	satisfied	050d30ffc32bea27eea9

Compiled by Annex 0.1.0 on 2026-09-15 from 34 source files in 69 ms. Rule packs: EU AI Act 2026.09.1 · GDPR (automated decisions) 2026.09.1 · NYC Local Law 144 2026.09.1 · NIST AI RMF 1.0 2026.09.1. This is a technical artefact, not legal advice, and not a conformity assessment.