

Technical documentation referred to in Article 11(1)

Regulation (EU) 2024/1689, as amended by Regulation (EU) 2026/1744, Annex IV · compiled from source by Annex

System	hireflow
Provider	hireflow
Version under assessment	dca401426407
Issued	2026-09-16
Classification	Contains a practice prohibited by Article 5: emotion inference in the workplace or an education setting.
Role under Articles 3(3) and 3(4)	provider and deployer
Evidence ledger	D94B - 458C - 8507 - 2B4B
Evidence citations	12
Open items	15

CONTENTS

1. General description of the AI system
2. Detailed description of the elements of the AI system and of the process for its development
3. Monitoring, functioning and control of the AI system
4. Appropriateness of the performance metrics
5. Risk management system in accordance with Article 9
6. Relevant changes made through the lifecycle
7. Harmonised standards applied, or solutions adopted
8. Copy of the EU declaration of conformity (Article 47)
9. Post-market performance evaluation system (Article 72)

§ Statement of method

This document was compiled by Annex directly from the source code of the system it describes, at the commit identified above. Every statement below falls into one of three categories, and they are never mixed:

1. **Evidenced.** The statement is supported by one or more citations to a file and a line in the repository. Each citation carries the SHA-256 digest of the file it came from, recorded in Appendix B.
2. **No evidence found.** The engine searched the whole snapshot and found nothing supporting the statement. The negative finding is recorded rather than omitted.

3. **Open.** The determination requires a judgement that cannot be derived from source code — a residual-risk acceptance, a declared accuracy level, an accountable person. These are listed per section and summarised on the front page. They are left blank deliberately. A generated document that invents them would be a false statement about the system, and where such a statement is supplied to a notified body or a national competent authority in reply to a request, Article 99(5) prices it at EUR 7 500 000 or 1 % of worldwide annual turnover.

No language model participated in any determination in this document. Classification and control evaluation are rule-based, deterministic and offline; the same commit always produces the same evidence ledger root. Where a language model contributed narrative prose, the section says so.

This document is a technical artefact, not legal advice, and it is not a conformity assessment. It is the input a competent person needs in order to carry one out, and the means by which a third party can check that its claims are true.

1 General description of the AI system

(a) Intended purpose, provider and version. Not determinable from the codebase: the operator did not state an intended purpose. This single sentence determines the risk classification under Article 6, so it cannot be left blank in a final version.

Provider: hireflow. Version under assessment: `dca401426407` .

(b) Interaction with other systems. The system calls externally provided models: OpenAI model provider. Article 25(1)(c) should be considered: any distributor, importer, deployer or other third party who modifies the intended purpose of a general-purpose AI system that has not been classified as high-risk and has already been placed on the market, such that it becomes high-risk under Article 6, becomes its provider. Whether those three conditions are met here is a determination for the operator; Annex records the model dependency, not the conclusion.

(c) Software versions. 7 production dependencies were resolved from the manifests in this repository. The full inventory is emitted as a CycloneDX ML-BOM alongside this document.

(d) Forms in which the system is placed on the market. Delivered as a hosted software service with a user-facing interface.

(e) Hardware. Not determinable from the codebase. Describe the runtime environment and any hardware requirement.

(f) Product photographs and layout. Not applicable: the system is software, not a component of a physical product.

(g)-(h) Deployer interface and instructions for use. No instructions for use aimed at a deployer were found.

EVIDENCE

```
src/lib/openai.ts:10
export const SCORING_MODEL = 'gpt-4o-mini';

src/lib/openai.ts:11
export const COPILOT_MODEL = 'gpt-4o';

src/app/api/chat/route.ts:9
const stream = await openai.chat.completions.create({

src/interview/signal.ts:18
const result = await openai.chat.completions.create({

src/screening/rank.ts:18
const completion = await openai.chat.completions.create({

tests/rank.test.ts:7
SCORING_MODEL: 'gpt-4o-mini',
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Intended purpose (point 1(a))

- ☐ Hardware and runtime environment (point 1(e))

2 Detailed description of the elements of the AI system and of the process for its development

(a) Development methods and recourse to third-party systems. The system is built on pre-trained models supplied by third parties (OpenAI model provider). Article 25(2) obliges the initial provider to provide the information and the reasonably expected technical access and other assistance needed to assess Article 16 compliance — unless it has specified that its system is not to be changed into a high-risk AI system, which switches that duty off. Confirm which applies before relying on it.

(b) Design specifications. The general logic is expressed in system prompts and orchestration code located in the evidence below. What the system optimises for, and the trade-offs accepted to meet Chapter III Section 2, are design decisions that must be recorded here. Annex IV point 2(b) also requires the **main design choices including rationale and assumptions made with regard to the persons or groups of persons on whom the system is intended to be used**. That half of the sub-point is the one an assessor probes hardest on an Annex III point 4(a) system, and it cannot be derived from source: state who the system is used on, and what was assumed about them.

(c) System architecture and computational resources. 14 source files were analysed. Not determinable from the codebase: describe how components feed into each other and the resources used to develop, train, test and validate the system.

(d) Data requirements and datasheets. No dataset documentation was found.

(e) Human oversight measures (Article 14). This system applies automated outcomes to people with no human review, no override path and no stop control.

(f) Pre-determined changes. Not determinable from the codebase. Record any changes to the system and its performance fixed at the time of the initial conformity assessment.

(g) Validation and testing procedures. No model evaluation suite was found for a high-risk system. Annex IV point 2(g) also requires metrics for potentially discriminatory impacts, and dated, signed test reports.

(h) Cybersecurity measures. No AI-specific security measures were found.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ General logic, key design choices and optimisation target (point 2(b))
- ☐ Design assumptions about the persons or groups the system is used on (point 2(b))
- ☐ System architecture and computational resources (point 2(c))
- ☐ Pre-determined changes (point 2(f))
- ☐ Dated and signed test reports (point 2(g))

3 Monitoring, functioning and control of the AI system

Capabilities and limitations in performance. Not determinable from the codebase from source alone. Annex IV point 3 asks specifically for the degrees of accuracy **for specific persons or groups** on whom the system is used, not only the overall figure.

Foreseeable unintended outcomes and sources of risk. Classification identified: Emotion inference in the workplace or an education setting; Employment: recruitment and candidate selection; Biometrics: emotion recognition; Solely automated decision with legal or similarly significant effect; Direct interaction with natural persons; Generation of synthetic text. Each is a source of risk to fundamental rights that Article 9 must address.

Human oversight measures. This system applies automated outcomes to people with no human review, no override path and no stop control.

Input data specifications. Not determinable from the codebase. Specify the expected input format, range and quality.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Accuracy for specific persons or groups
- ☐ Input data specifications

4 Appropriateness of the performance metrics

Not determinable from the codebase: no evaluation suite was found, so there are no metrics whose appropriateness could be assessed.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Justification that the chosen metrics suit the intended purpose

5 Risk management system in accordance with Article 9

No evidence found — Risk management system across the lifecycle (EU AI Act Art. 9). No risk management documentation was found for a high-risk system.

Article 9(5) requires residual risk to be judged acceptable for each individual hazard and overall. That judgement, and the person who made it, cannot be derived from code.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Residual risk acceptance and the accountable person (Article 9(5))

6 Relevant changes made through the lifecycle

This document describes the system at commit `dca401426407` . Annex maintains change history by comparing consecutive scans: where a change affects compliance with Chapter III Section 2 or modifies the intended purpose, it is a substantial modification within Article 3(23) and reopens the conformity assessment under Article 43(4).

Not determinable from the codebase: no change control artefacts were found. Article 17(1)(a) expects management of modifications to be part of the quality management system.

EVIDENCE

```
tests/rank.test.ts:10
describe('rankApplicants', () => {
.....
tests/rank.test.ts:11
it('sorts advanced applicants by score', async () => {
.....
.github/workflows/test.yml:5
runs-on: ubuntu-latest
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Log of relevant changes since the initial conformity assessment

7 Harmonised standards applied, or solutions adopted

No harmonised standards for the AI Act have yet been cited in the Official Journal of the European Union. CEN-CENELEC announced EN 18286 — a quality management system supporting AI Act regulatory purposes — as the first European standard in this family on 30 July 2026, but a standard is only harmonised once its reference is published in the OJEU, and until then Article 40 gives no presumption of conformity. So Annex IV point 7 takes its second branch: a detailed description of the solutions adopted to meet the Chapter III Section 2 requirements, and a list of other relevant standards applied.

No other standards or technical specifications were referenced in this repository.

The solutions adopted are the controls assessed in this document; each is listed with its evidence in the appendix.

EVIDENCE

```
tests/rank.test.ts:10
describe('rankApplicants', () => {
```

```
tests/rank.test.ts:11
it('sorts advanced applicants by score', async () => {
```

```
.....
.github/workflows/test.yml:5
runs-on: ubuntu-latest
```

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ List of other standards and technical specifications applied

8 Copy of the EU declaration of conformity (Article 47)

Article 47 requires a written, machine-readable, physically or electronically signed EU declaration of conformity for each high-risk AI system, kept available to national competent authorities for ten years.

Not determinable from the codebase: no EU declaration of conformity was found. It cannot be drawn up until the open items in this document are closed — by drawing it up, the provider assumes responsibility for compliance.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ EU declaration of conformity (Article 47, Annex V)

9 Post-market performance evaluation system (Article 72)

No evidence found — Post-market monitoring plan (EU AI Act Art. 72(1)). Neither a post-market monitoring plan nor production monitoring was found.

Article 72(3) makes the post-market monitoring plan part of this technical documentation. The Commission template is not due until 2 September 2027, so the plan is free-form until then.

OPEN ITEMS — DETERMINATIONS THIS DOCUMENT CANNOT MAKE

- ☐ Post-market monitoring plan (Article 72(3))

A Appendix A — Control register

Control	Citation	Status	Finding
No emotion inference in the workplace or education	EU AI Act Art. 5(1)(f)	MISSING	This repository infers emotional or affective state from people in a recruitment, employment or education context. Article 5(1)(f) prohibits that outright.
Measures to support AI literacy	EU AI Act Art. 4	MISSING	No AI literacy or AI usage policy documentation was found.
Tell people they are talking to an AI	EU AI Act Art. 50(1)	MISSING	This system talks directly to people and no AI disclosure was found anywhere in the codebase.
Mark synthetic output in a machine-readable format	EU AI Act Art. 50(2)	MISSING	This system generates synthetic content and no machine-readable marking was found.
Notify people exposed to emotion recognition or biometric categorisation	EU AI Act Art. 50(3)	MISSING	An emotion recognition or biometric categorisation capability was detected with no notification to the people exposed to it.
Determine and record whether you are the provider or the deployer	EU AI Act Art. 3(3)	NEEDS REVIEW	This repository calls a third-party model and builds a product around the result. If that product is placed on the market or put into service under your own name or trademark, Article 3(3) makes you its provider rather than merely a deployer — and Annex cannot

Control	Citation	Status	Finding
			settle that from source code alone.
Risk management system across the lifecycle	EU AI Act Art. 9	MISSING	No risk management documentation was found for a high-risk system.
Examine training and evaluation data for bias	EU AI Act Art. 10(2)(f)	MISSING	No bias or fairness testing was found for a high-risk system. Protected attributes were found in the data model, which raises the stakes: they can be used to measure disparity, or to cause it.
Document data provenance and preparation	EU AI Act Art. 10(2)	MISSING	No dataset documentation was found.
Annex IV technical documentation exists and is current	EU AI Act Art. 11(1)	MISSING	No Annex IV technical documentation was found.
Automatic logging of events over the system lifetime	EU AI Act Art. 12(1)	MISSING	Structured application logging exists, but nothing records the model inputs, outputs and versions behind each decision. Application logs are not Article 12 logs.
Instructions for use for deployers	EU AI Act Art. 13(3)	MISSING	No instructions for use aimed at a deployer were found.
Effective human oversight while the system is in use	EU AI Act Art. 14(1)	MISSING	This system applies automated outcomes to people with no human review, no override path and no stop control.
Declared accuracy metrics measured against a fixed set	EU AI Act Art. 15(1)	MISSING	No model evaluation suite was found for a high-risk system.
AI-specific cybersecurity measures	EU AI Act Art. 15(5)	MISSING	No AI-specific security measures were found.
Quality management system	EU AI Act Art. 17(1)	PARTIAL	Found a CI pipeline and an automated test suite, which is a start on Article 17.

Control	Citation	Status	Finding
Post-market monitoring plan	EU AI Act Art. 72(1)	MISSING	Neither a post-market monitoring plan nor production monitoring was found.
Serious incident reporting procedure	EU AI Act Art. 73(2)	MISSING	No serious incident reporting procedure was found.
Registration in the EU database	EU AI Act Art. 49(1)	MISSING	No reference to EU database registration was found.
Deployer duties: assigned overseers, input data, and telling the workforce	EU AI Act Art. 26(2)	MISSING	This system makes decisions about people at work, and nothing records that workers or their representatives were informed before it was put into service.
Explain an individual decision to the person it was taken about	EU AI Act Art. 86(1)	MISSING	Nothing was found that would answer an affected person asking why a decision about them came out the way it did.
Right to human intervention in an automated decision	GDPR Art. 22(1)	MISSING	An automated decision is taken about people with no human intervention path and no way to contest it.
Meaningful information about the logic involved	GDPR Art. 13(2)(f)	MISSING	No explanation capability and no privacy notice were found.
Data protection impact assessment	GDPR Art. 35(1)	MISSING	No data protection impact assessment was found.
Erasure and rectification of personal data	GDPR Art. 16	MISSING	No erasure or rectification path was found for personal data.
Lawful basis for special-category data	GDPR Art. 9(1)	MISSING	Special-category personal data appears in the data model with no recorded lawful basis.
Annual independent bias audit	NYC Local Law 144 § 5-301(a)	MISSING	This repository screens candidates or employees and no bias audit was found.
Selection rate and impact ratio by category	NYC Local Law 144 § 5-301(b)	MISSING	No selection-rate or impact-ratio calculation was found.
Publish the audit summary before use	NYC Local Law 144 § 5-303	MISSING	No published bias audit summary was found.

Control	Citation	Status	Finding
Ten business days notice to candidates	NYC Local Law 144 § 5-304	MISSING	No candidate notice was found for an automated employment decision tool.
Publish the data retention policy and data sources	NYC Local Law 144 6 RCNY § 5-304(d)	MISSING	Neither a data retention policy nor data source documentation was found.
Clear pre-use notice before a consequential decision	Colorado ADMT Act § 6-1-1703	MISSING	No pre-use notice was found ahead of a consequential decision.
Explain an adverse outcome within thirty days	Colorado ADMT Act § 6-1-1703	MISSING	No adverse-outcome explanation path was found.
Designated reviewers authorised to override	Colorado ADMT Act § 6-1-1704	MISSING	No human review or override capability was found for a consequential decision.
Right to see and correct the data used	Colorado ADMT Act § 6-1-1705	MISSING	No way for a consumer to see or correct the data used in the decision was found.
Documentation handed to deployers	Colorado ADMT Act § 6-1-1702	MISSING	No documentation aimed at a downstream deployer was found.
GOVERN 1.1 — legal and regulatory requirements are understood and documented	NIST AI RMF GOVERN 1.1	MISSING	No documented understanding of the regulatory requirements that apply to this system was found.
GOVERN 1.6 — an inventory of AI systems is maintained	NIST AI RMF GOVERN 1.6	MISSING	No AI system or model inventory was found.
MAP 1.1 — intended purpose and deployment context are documented	NIST AI RMF MAP 1.1	MISSING	The intended purpose of the system is not documented anywhere in the repository.
MEASURE 2.11 — fairness and bias are evaluated and documented	NIST AI RMF MEASURE 2.11	MISSING	No fairness or bias evaluation was found.

Control	Citation	Status	Finding
MEASURE 2.7 — security and resilience are evaluated	NIST AI RMF MEASURE 2.7	MISSING	No AI-specific security evaluation was found.
MANAGE 4.1 — post-deployment monitoring plans are implemented	NIST AI RMF MANAGE 4.1	PARTIAL	Found change management.

B Appendix B — Evidence ledger

Algorithm `sha256-chain/v1`. Root
`d94b458c85072b4bf5d31c62a27da6301eaec37adb5b0b8104e58c121100878d`.

Each entry hashes the previous entry, the control identifier, its status and score, the rule-pack version, and the digest of every piece of evidence cited. Re-running Annex on the same commit reproduces this root exactly. `annex verify` re-derives every entry from the results this document describes, so an edited status no longer hashes to its recorded value; `annex verify --against <dir>` additionally re-hashes each cited file off disk, so a document that no longer describes the tree says so. This is a checksum chain, not a signature: it makes a silent edit detectable by anyone holding the source, not impossible.

#	Control	Status	Entry hash
0	colorado-admt.adverse-explanation	missing	<code>ba0e9aa1c37658515bbf</code>
1	colorado-admt.data-correction	missing	<code>a0b8b917a853f38c3f83</code>
2	colorado-admt.developer-documentation	missing	<code>ee6d5a7740d1a657ea4b</code>
3	colorado-admt.human-review-override	missing	<code>a82878e0bc3696cf18c5</code>
4	colorado-admt.pre-use-notice	missing	<code>72c01e4694404a22fb10</code>
5	eu-ai-act.art10.bias-examination	missing	<code>bd7189b7eb9cf38d6120</code>
6	eu-ai-act.art10.data-governance	missing	<code>13cfe95ee2188080fa99</code>
7	eu-ai-act.art11.technical-documentation	missing	<code>d1900f4f5a2520467eec</code>
8	eu-ai-act.art12.record-keeping	missing	<code>ee25edeb22836054d3e7</code>
9	eu-ai-act.art13.instructions-for-use	missing	<code>b382f6740e17b473f2b7</code>
10	eu-ai-act.art14.human-oversight	missing	<code>60bc3d353ee1ca1bd487</code>
11	eu-ai-act.art15.accuracy	missing	<code>24bc066e20726d7bdcca</code>
12	eu-ai-act.art15.cybersecurity	missing	<code>b82b0ced6a71f7071ed0</code>

#	Control	Status	Entry hash
13	eu-ai-act.art17.quality-management	partial	3e0e620f0816fb736364
14	eu-ai-act.art19.log-retention	not applicable	d15582cf70501b898332
15	eu-ai-act.art26.deployer-obligations	missing	ff15420ab333f2e922be
16	eu-ai-act.art27.fria	not applicable	a2733ae175d3c9c11ca9
17	eu-ai-act.art3.role-determination	needs review	a4d7d1238422d8609231
18	eu-ai-act.art4.ai-literacy	missing	72420d48474dd532c123
19	eu-ai-act.art49.registration	missing	c3634fc436af16f079ce
20	eu-ai-act.art5.emotion-workplace	missing	a0a6c0f217dc82258426
21	eu-ai-act.art5.face-scraping	not applicable	aecff6ea8aa77ff8b039
22	eu-ai-act.art5.ncii-csam-safeguards	not applicable	244cd46aaf8658c84b97
23	eu-ai-act.art5.social-scoring	not applicable	ace868aa00803b239a36
24	eu-ai-act.art50.1.interaction-disclosure	missing	8178a31a36f950d99d74
25	eu-ai-act.art50.2.content-marking	missing	bd39c80d7f9a7075cdd0
26	eu-ai-act.art50.3.biometric-notification	missing	95c4289a1cc9e315c1d5
27	eu-ai-act.art50.4.deepfake-labelling	not applicable	85743a23c9bf98ec9017
28	eu-ai-act.art72.post-market-monitoring	missing	7862c6e8228e0d2429d4
29	eu-ai-act.art73.incident-reporting	missing	3a6f2eac1caa18afee6e
30	eu-ai-act.art86.right-to-explanation	missing	0f16ae00f83e4512ef51
31	eu-ai-act.art9.risk-management	missing	5da84a6503bac978cc26
32	gdpr.art13-15.meaningful-information	missing	aa9c30477fb26f48ae06
33	gdpr.art17.erasure	missing	925e259d537dc98c04f1
34	gdpr.art22.human-intervention	missing	6cb1e3d445445e4840c7
35	gdpr.art35.dpia	missing	81105d16cfb8c3b5b42a
36	gdpr.art9.special-category	missing	32757a4fe96d3a8f54e0
37	nist-ai-rmf.govern-1-1	missing	601ff5a8a0069abb1c2a
38	nist-ai-rmf.govern-1-6	missing	433de1d8071a27cf66de
39	nist-ai-rmf.manage-4-1	partial	64023eabbff38fa3b54a

#	Control	Status	Entry hash
40	nist-ai-rmf.map-1-1	missing	38259bda13f06d815832
41	nist-ai-rmf.measure-2-11	missing	b1850c48d4bfc6dd0194
42	nist-ai-rmf.measure-2-7	missing	280bc18870a0f88d7303
43	nyc-ll144.bias-audit	missing	3ddd7baf877a82798f83
44	nyc-ll144.candidate-notice	missing	30b579754fcf859ec4f7
45	nyc-ll144.data-policy	missing	0a4fba6813afc0ddda07
46	nyc-ll144.impact-ratio	missing	3dbcabc28ec3e01a4622
47	nyc-ll144.publish-summary	missing	d94b458c85072b4bf5d3

Compiled by Annex 0.1.0 on 2026-09-16 from 14 source files in 50 ms. Rule packs: EU AI Act 2026.09.1 · GDPR (automated decisions) 2026.09.1 · NYC Local Law 144 2026.09.1 · Colorado ADMT Act 2026.09.1 · NIST AI RMF 1.0 2026.09.1. This is a technical artefact, not legal advice, and not a conformity assessment.